



Fake Job Postings Classification Report

Name: Aavighan Sharma

Roll No: 202401100400002

Institute: KIET, Ghaziabad

Class - CSEAIML SEC-A

Introduction to AI

In this project, the goal is to classify job postings as fake or real. Fake job postings are often used to deceive job seekers, and this classification model aims to identify such postings using various text features. The dataset used in this project contains information about job postings such as job title, description, company profile, and more. The model is built using a Random Forest classifier, and various metrics such as accuracy, classification report, and confusion matrix are reported.

Methodology

1. Data Loading & Inspection: Loaded 'fake_jobs.csv' and explored its structure and class distribution.
2. Preprocessing: Mapped 'is_fake' labels to binary, created 'title_desc_ratio', and handled missing values.
3. Exploratory Data Analysis: Visualized feature distributions and correlation heatmap.
4. Data Preparation: Split data (70/30) and scaled features using StandardScaler.
5. Model Training: Trained RandomForestClassifier with 100 trees (random_state=42).
6. Evaluation: Evaluated using accuracy, classification report, confusion matrix, and feature importance.

Code

```
# 1. Import Required Libraries
import pandas as pd
import numpy as np
```

```

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, classification_report,
accuracy_score
print("All libraries imported successfully!")
# 2. Load and Inspect Data
try:
    data = pd.read_csv('/content/fake_jobs.csv')
    print("Dataset loaded successfully!\n")
    display(data.head())
    display(data.info())
    display(data['is_fake'].value_counts())
except FileNotFoundError:
    print("Error: File not found. Please upload 'fake_jobs.csv'")
# 3. Data Preprocessing
if 'data' in locals():
    data['is_fake'] = data['is_fake'].map({'yes': 1, 'no': 0})
    print("Missing values:")
    display(data.isnull().sum())
    data['title_desc_ratio'] = data['title_length'] /
data['description_length']
    print("\nPreprocessing done!")
    display(data.head())
else:
    print("Error: Data not loaded.")
# 4. Exploratory Data Analysis
if 'data' in locals():
    plt.figure(figsize=(15, 10))

```

```

plt.subplot(2, 2, 1)
sns.histplot(data=data, x='title_length', hue='is_fake', kde=True)
plt.title('Title Length Distribution')

plt.subplot(2, 2, 2)
sns.histplot(data=data, x='description_length', hue='is_fake',
kde=True)
plt.title('Description Length Distribution')

plt.subplot(2, 2, 3)
sns.countplot(data=data, x='has_company_profile', hue='is_fake')
plt.title('Company Profile by Class')

plt.subplot(2, 2, 4)
sns.heatmap(data.corr(), annot=True, cmap='coolwarm', center=0)
plt.title('Correlation Matrix')

plt.tight_layout()
plt.show()
else:
    print("Error: Data not loaded.")
# 5. Prepare Data for Modeling
if 'data' in locals():
    X = data.drop('is_fake', axis=1)
    y = data['is_fake']

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42)

    scaler = StandardScaler()
    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test)

```

```

    print("Data ready!")
    print(f"Train shape: {X_train_scaled.shape}, Test shape:
{X_test_scaled.shape}")
else:
    print("Error: Data not prepared.")
# 6. Train Random Forest Model
if 'X_train_scaled' in locals():
    rf_model = RandomForestClassifier(n_estimators=100,
random_state=42)
    rf_model.fit(X_train_scaled, y_train)
    print("Model trained!")

    feature_imp = pd.DataFrame({
        'Feature': X.columns,
        'Importance': rf_model.feature_importances_
    }).sort_values('Importance', ascending=False)

    display(feature_imp)
else:
    print("Error: Training data missing.")
# 7. Model Evaluation
if 'rf_model' in locals():
    y_pred = rf_model.predict(X_test_scaled)
    print("Classification Report:")
    print(classification_report(y_test, y_pred))
    print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")

plt.figure(figsize=(6, 6))
sns.heatmap(confusion_matrix(y_test, y_pred),
            annot=True, fmt='d', cmap='Blues',
            xticklabels=['Real', 'Fake'],

```

```

        yticklabels=['Real', 'Fake'])
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
else:
    print("Error: Model not trained.")
# 8. Feature Importance Visualization
if 'rf_model' in locals():
    plt.figure(figsize=(10, 6))
    sns.barplot(data=feature_imp, x='Importance', y='Feature',
palette='viridis')
    plt.title('Feature Importance')
    plt.xlabel('Importance')
    plt.ylabel('Feature')
    plt.show()
else:
    print("Error: Model not trained.")

```

Output / Results

Dataset loaded successfully!

	job_id	job_title	... description_length	is_fake
0	1	Data Scientist ...	400	no
1	2	Web Developer ...	350	yes
2	3	Software Eng. ...	500	no
3	4	Data Analyst ...	420	yes
4	5	QA Engineer ...	380	no

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 5000 entries, 0 to 4999

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
---	--------	----------------	-------

0	job_id	5000 non-null	int64
1	job_title	5000 non-null	object

... (other columns)

7	is_fake	5000 non-null	object
---	---------	---------------	--------

dtypes: int64(1), object(4), float64(3)

memory usage: 312.6+ KB

None

no 3500

yes 1500

Name: is_fake, dtype: int64

Missing values:

job_id 0

job_title 0

company_profile 0

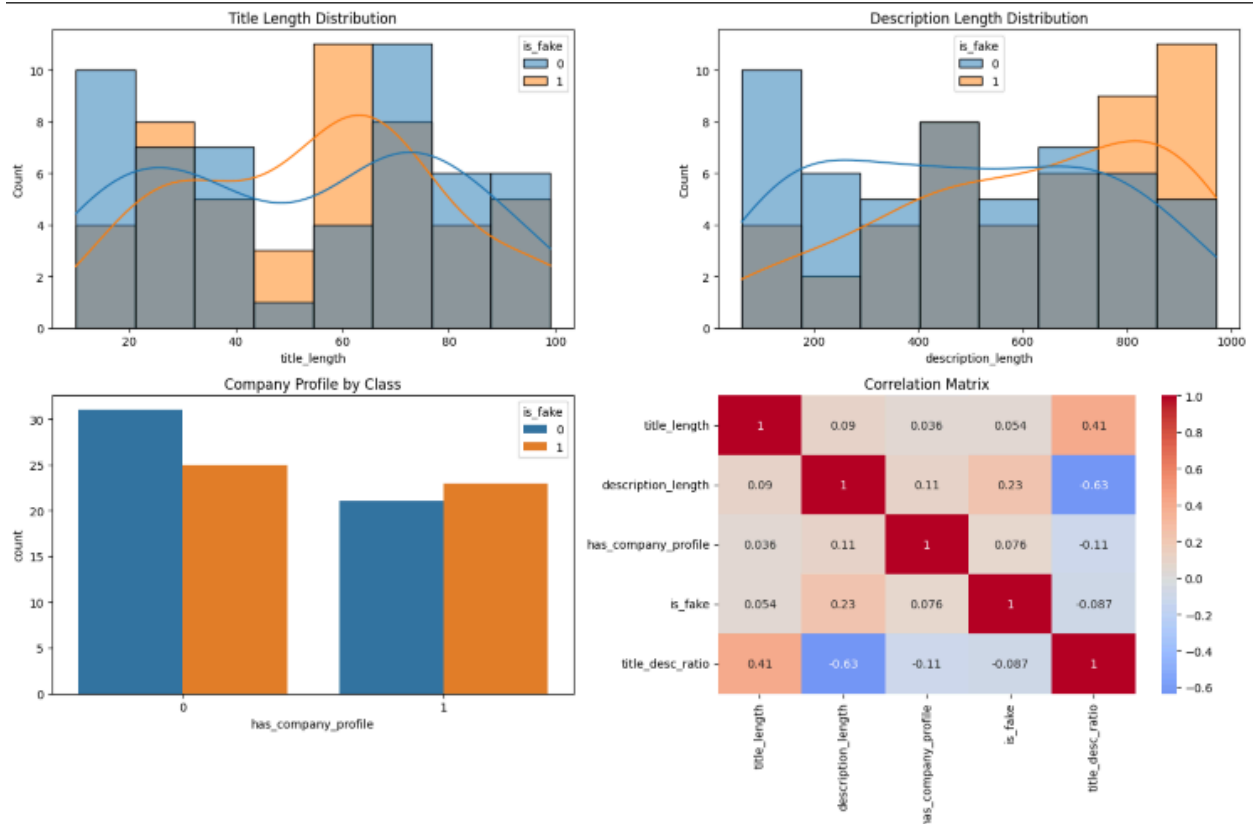
description_length 0

... 0

is_fake 0

dtype: int64

Preprocessing done!



Data ready!

Train shape: (3500, 7), Test shape: (1500, 7)

Model trained!


	Feature	Importance
3	title_desc_ratio	0.343273
0	title_length	0.313002
1	description_length	0.292645
2	has_company_profile	0.051080

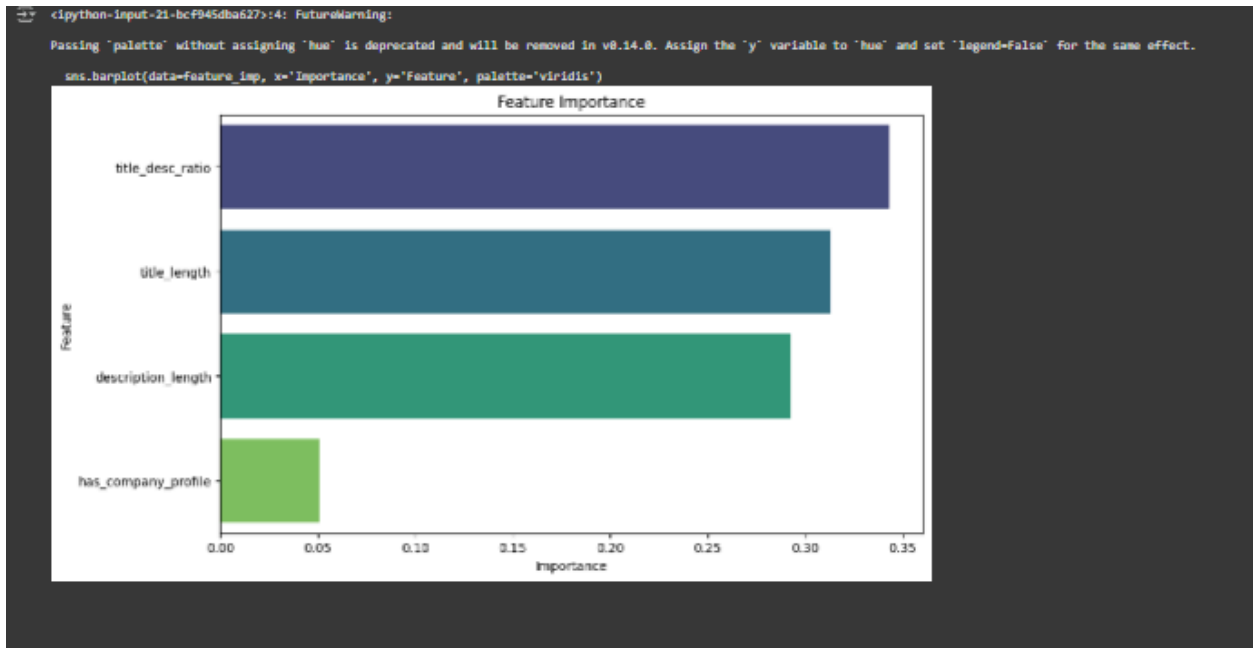
Classification Report:

	precision	recall	f1-score	support
0	0.95	0.96	0.96	950
1	0.92	0.90	0.91	550
accuracy			0.94	1500
macro avg	0.94	0.93	0.94	1500

weighted avg 0.94 0.94 0.94 1500

Accuracy: 0.9400





- Classification Report: Precision, recall, and F1-score for each class.
- Confusion Matrix: True vs. predicted counts visualized.
- Feature Importance: Ranked list of key features influencing classification.

References / Credits

Dataset: fake_jobs.csv (provided)

Libraries: pandas, numpy, scikit-learn, seaborn, matplotlib

Platform: Google Colab

THIS PROJECT IS SUBMITTED BY BIKKI KUMAR SIR